


k-Nearest Neighbors & Decision Trees

Outline

- k-Nearest Neighbors (kNN)
 - Algorithm
 - Distance metrics
 - kNN tradeoffs
- Decision Trees
 - How to build a decision tree
 - Information gain vs. gini impurity
 - Decision tree tradeoffs
 - Pruning

k-Nearest Neighbors (kNN)



Distance Metrics

- Euclidean distance

$$\text{dist}(a, b) = \|a - b\| = \sqrt{(a - b) \cdot (a - b)} = \sqrt{\sum_i (a_i - b_i)^2}$$

- Cosine similarity

$$\text{dist}(a, b) = \frac{a \cdot b}{\|a\| \|b\|} = \frac{\sum_i a_i b_i}{\sqrt{\sum_i a_i^2} \sqrt{\sum_i b_i^2}}$$

Pseudocode - kNN

For datapoint in training set:

- calculate distance from datapoint to new_value

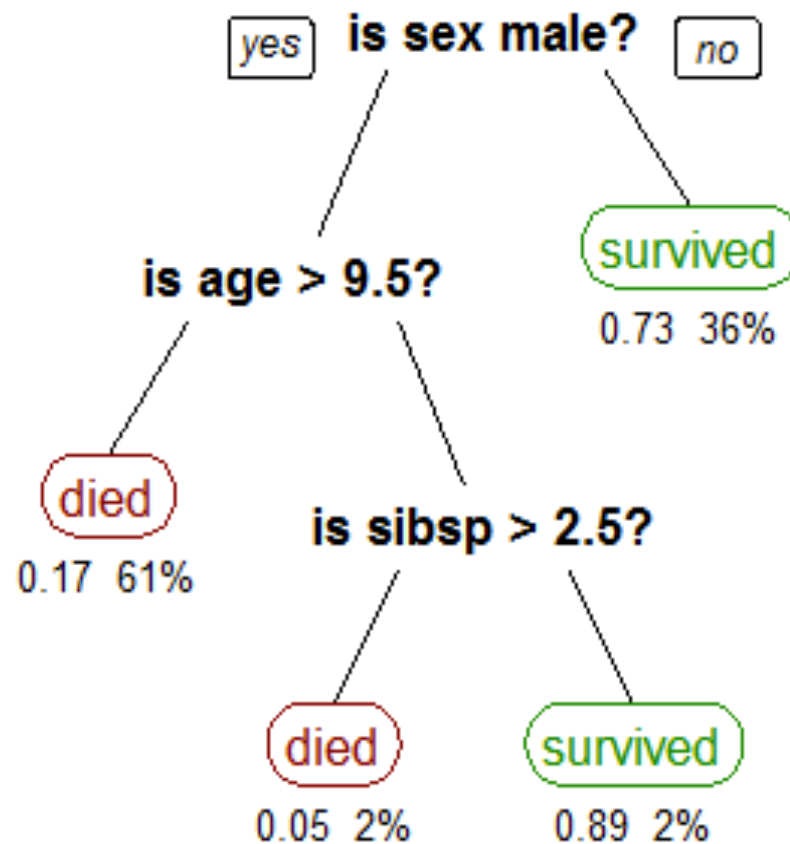
Order distances in increasing order and take the first k

Take the label with the most votes

kNN Tradeoffs

- Pros
 - Really easy to train (just save all the data)
 - Easily works with any number of classes
 - Easy to add new training datapoints
- Con
 - Really slow to predict (especially if you have a lot of features)

Decision Trees



How to Build a Decision Tree

- To do a binary split:
 - For a categorical variable, choose either value or not value (e.g. sunny or not sunny)
 - For a continuous variable, choose a threshold and do $>$ or \leq the value (e.g. temperature <75 or ≥ 75)
- To measure how good the split is:
 - Information gain
 - Gini impurity


Entropy

- Entropy – a measure of the amount of disorder in a set

$$H(y) = - \sum_{i=1}^m P(c_i) \log_2 (P(c_i))$$

$P(c)$ is the percent of the group that belongs to a given class

More on Entropy



Information Gain

$$Gain(S, D) = H(S) - \sum_{V \in D} \frac{|V|}{|S|} H(V)$$

S is the original set and D is the splitting of the set (a partition), each V is a subset of S

Gini Impurity

- It is a measure of this probability:
 - Take random element from the set
 - Label it randomly according to the distribution of labels in the set
 - What is the probability that is it labeled incorrectly?

$$Gini = \sum_{i=1}^m P(c_i)(1 - P(c_i)) = 1 - \sum_{i=1}^m P(c_i)^2$$

Pseudocode – Decision Trees

function BuildTree:

- If every item in the dataset is in the same class or there is no feature left to split the data:

 - return a leaf node with the class label

- Else:

 - find the best feature and value to split the data

 - split the dataset

 - create a node

 - for each split

 - call BuildTree and add the result as a child of the node

 - return node

Decision Tree Tradeoffs

- Pros
 - Easily interpretable
 - Handles missing values and outliers
 - Non-parametric/non-linear/model complex phenomenon
 - Computationally cheap to predict
 - Can handle irrelevant features
 - Mixed data (discrete and continuous)
- Cons
 - Computationally expensive to train
 - Greedy algorithm (local optima)
 - Very easy to overfit

Pruning - Prepruning

- Making the decision tree algorithm stop early
 - Leaf size: stop when the number of data points for a leaf gets below a threshold
 - Depth: stop when the depth of the tree (distance from root to leaf) reaches a threshold
 - Mostly the same: stop when some percent of the data points are the same (rather than all the same)
 - Error threshold: stop when the error reduction (information gain) is not improved significantly

Pruning - Postpruning

- Involves building the tree first and then choosing to cut off some of the leaves
- Psuedocode:

function Prune:

 if either left or right is not a leaf:

 call Prune on that split

 if both left and right are leaf nodes:

 calculate error associated with merging two nodes

 calculate error associated without merging two nodes

 if merging results in lower error:

 merge the leaf nodes

Decision Tree Variants

- ID3 (Iterative Dichotomiser 3)
 - Designed for only categorical features
 - Splits categorical features completely
 - Uses entropy and information gain to pick the best split
- CART (Classification and Regression Tree)
 - Handles both categorical and continuous data
 - Always uses binary splits
 - Uses gini impurity to pick the best split
- C4.5
 - Handles continuous data
 - Implements pruning to reduce overfitting
- C5.0
 - Proprietary, do not have access to the specifics

In Practice

- We always implement pruning to avoid overfitting
- Either gini or information gain is acceptable
- Sometimes fully splitting categorical features is preferred, but generally we air on the side of binary splits

sklearn

- Pruning with `max_depth`, `min_samples_split`, `min_samples_leaf` or `max_leaf_nodes`
- Gini is default, but you can also choose entropy
- Does binary splits (you would need to binarize categorical features)

Regression Trees

- Decision trees can also be used for regression
- Can average the values at each leaf node to predict a continuous value
- Can also use a combination of decision trees and linear regression on the leaf nodes (model trees)